


File I/O

Zilogic Systems

1. Input and Output

- Each process has a number of input and output sources associated with it.
- Examples:
 - Input Source: data from a file
 - Output Source: data sent to printer
 - Input Source: input from keyboard
 - Output Source: message printed to screen
- Each input/output source is given a unique number within the process.
- There are three input/output sources available by default to a process.
 - 0 - standard input, input source, usually associated with the terminal
 - 1 - standard output, output source, usually associated with the terminal
 - 2 - standard error, output source, also associated with the terminal
- The unique numbers give to each source is called a file descriptor.

Figure 1. Standard Input, Output and Error



- Redirecting input, output and error.

Figure 2. Redirecting Standard Input





Figure 3. Redirecting Standard Output

2. Read and Write

- Data can be got from input sources using `read()` system call.
- Data can be sent to output sources using `write()` system call.

```
#include <unistd.h>

ret = read(fd, buf, len)

int fd
char *buf
size_t len
ssize_t ret
```

- `fd` specifies the file descriptor to read from
- `buf` is a buffer (char array) into which read data will be stored
- `len` is the size of the buffer
- returns the no. of bytes read
- if `ret` is `0`, indicates end-of-file
- if `ret` is `-1`, indicates error has occurred

```
#include <unistd.h>

ret = write(fd, buf, len)

int fd
char *buf
size_t len
ssize_t ret
```

- `fd` specifies the file descriptor to write to
- `buf` is a buffer from which data will be written
- `len` is the no. of bytes in the buffer
- returns the no. of bytes written
- if `ret` is `-1`, indicates error has occurred

Char array vs. String

- The buffer passed to `read()` and `write()` is an character array not a string.

3. Implementing `tr`

`tr` initial version.

```
#include <unistd.h>

#define BUF_SIZE 256

int main(int argc, char *argv[])
{
    char buf[BUF_SIZE];
    ssize_t rlen;
    int i;
    char from;
    char to;

    from = 'e';
    to = 'a';

    while (1) {
        rlen = read(0, buf, sizeof(buf));
        if (rlen == 0)
            return 0;

        for (i = 0; i < rlen; i++) {
            if (buf[i] == from)
                buf[i] = to;
        }

        write(1, buf, rlen);
    }

    return 0;
}
```

`tr` with error checking.

```
#include <unistd.h>
#include <error.h>
#include <errno.h>

#define BUF_SIZE 256

int main(int argc, char *argv[])
{
    char buf[BUF_SIZE];
    ssize_t rlen, wlen;
    int i;
    char from;
    char to;
```

```

    from = 'e';
    to = 'a';

    while (1) {
        rlen = read(0, buf, sizeof(buf));
        if (rlen == -1) /* ❶ */
            error(1, errno, "error reading input"); /* ❷ */
        if (rlen == 0)
            return 0;

        for (i = 0; i < rlen; i++) {
            if (buf[i] == from)
                buf[i] = to;
        }


        wlen = write(1, buf, rlen);
        if (wlen == -1) /* ❸ */
            error(1, errno, "error writing to output");
    }

    return 0;
}

```

- ❶ `read()` returns -1 on error. Check for errors and terminate.
- ❷ `error()` prints an error message and terminates program.
- ❸ `write()` returns -1 on error. Check for errors and terminate.

Figure 4. Handling Partial Writes



tr with partial write handling.

```

#include <unistd.h>
#include <error.h>
#include <errno.h>

#define BUF_SIZE 256

int main(int argc, char *argv[])
{
    char buf[BUF_SIZE];
    char *bufp;
    ssize_t rlen, wlen;
    int i;
    char from;
    char to;

```

```

    from = 'e';
    to = 'a';

    while (1) {
        rlen = read(0, buf, sizeof(buf));
        if (rlen == -1)
            error(1, errno, "error reading input");
        if (rlen == 0)
            return 0;

        for (i = 0; i < rlen; i++) {
            if (buf[i] == from)
                buf[i] = to;
        }

        wlen = 0;
        bufp = buf;
        do { /* ❶ */
            wlen = write(1, bufp, rlen);
            if (wlen == -1)
                error(1, errno, "error writing to output");

            rlen -= wlen;
            bufp += wlen;
        } while (rlen != 0); /* ❷ */
    }

    return 0;
}

```

❶❷ Loop till all data has been written.

tr with macro names for fds.

```

#include <unistd.h>
#include <error.h>
#include <errno.h>

#define BUF_SIZE 256

int main(int argc, char *argv[])
{
    char buf[BUF_SIZE];
    char *bufp;
    ssize_t rlen, wlen;
    int i;
    char from;
    char to;

    from = 'e';
    to = 'a';

    while (1) {

```

```
        rlen = read(STDIN_FILENO, buf, sizeof(buf));
        if (rlen == -1)
            error(1, errno, "error reading input");
        if (rlen == 0)
            return 0;

        for (i = 0; i < rlen; i++) {
            if (buf[i] == from)
                buf[i] = to;
        }

        wlen = 0;
        bufp = buf;
        do {
            wlen = write(STDOUT_FILENO, bufp, rlen);
            if (wlen == -1)
                error(1, errno, "error writing to output");

            rlen -= wlen;
            bufp += wlen;
        } while (rlen != 0);
    }

    return 0;
}
```

tr with arguments.

```
#include <unistd.h>
#include <error.h>
#include <errno.h>

#define BUF_SIZE 256

int main(int argc, char *argv[])
{
    char buf[BUF_SIZE];
    char *bufp;
    ssize_t rlen, wlen;
    int i;
    char from;
    char to;

    if (argc != 3) {
        error(1, 0, "too many args or required args not specified");
    }

    from = argv[1][0];
    to = argv[2][0];

    while (1) {
        rlen = read(STDIN_FILENO, buf, sizeof(buf));
        if (rlen == -1)
```

```
        error(1, errno, "error reading input");
    if (rlen == 0)
        return 0;

    for (i = 0; i < rlen; i++) {
        if (buf[i] == from)
            buf[i] = to;
    }

    wlen = 0;
    bufp = buf;
    do {
        wlen = write(STDOUT_FILENO, bufp, rlen);
        if (wlen == -1)
            error(1, errno, "error writing to output");

        rlen -= wlen;
        bufp += wlen;
    } while (rlen != 0);
}

return 0;
}
```

4. Open and Close

- Additional file descriptors can be obtained by opening files, sockets, device nodes, etc.

```
#include <fcntl.h>

fd = open(name, flags)

char *name
int flags
int fd
```

- `name` is the path name to the file.
- `flags` - `O_RDONLY`, `O_WRONLY`, `O_RDWR`
- returns a file descriptor on success
- returns `-1` on error

```
#include <fcntl.h>

ret = close(fd)

int fd
int ret
```

- `fd` is the file descriptor to be closed
- returns `0` on success and `-1` on error

5. Implementing `cat`

`cat` initial version.

```
#include <unistd.h>
#include <fcntl.h>

#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    char *filename;
    int fd;
    char buf[1024];
    ssize_t ret = 1;
    int i;

    if (argc != 2) {
        printf("Usage: mycat <filename>\n");
        exit(1);
    }

    filename = argv[1];

    fd = open(filename, O_RDONLY);
    while (ret != 0) {
        /* Reading into buffer */
        ret = read(fd, buf, sizeof(buf));
        for (i = 0; i < ret; i++) {
            putchar(buf[i]);
        }
    }

    close(fd);

    return 0;
}
```

`cat` with error checking.

```
#include <unistd.h>
#include <fcntl.h>

#include <stdio.h>
#include <stdlib.h>
#include <error.h>
#include <errno.h>

int main(int argc, char *argv[])
{
    char *filename;
    int fd;
    char buf[1024];
```



```
    ssize_t ret = 1;
    int i;

    if (argc != 2) {
        printf("Usage: mycat <filename>\n");
        exit(1);
    }

    filename = argv[1];

    fd = open(filename, O_RDONLY);
    if (fd == -1) { /* ❶ */
        error(1, errno, "error opening file %s", filename); /* ❷ */
    }
    while (ret != 0) {
        ret = read(fd, buf, sizeof(buf));
        if (ret == -1) { /* ❸ */
            error(1, errno, "error reading file %s", filename);
        }
        for (i = 0; i < ret; i++) {
            putchar(buf[i]);
        }
    }

    if (close(fd) == -1) {
        error(1, errno, "error closing file %s", filename);
    }

    return 0;
}
```

- ❶ Checking for error in `open()` syscall.
- ❷ Prints an error message and terminates the process.
- ❸ Checking for error in `read()` syscall.

6. Implementing `cp`

`cp` initial version.

```
#include <unistd.h>
#include <fcntl.h>

#include <stdio.h>
#include <stdlib.h>
#include <error.h>
#include <errno.h>

void write_buf(int fd, char *bufp, size_t nbytes, char *filename)
{
    int written;

    while (nbytes != 0) {
        written = write(fd, bufp, nbytes);
        if (written == -1)
```

```
        error(1, errno, "error writing to file %s", filename);

        if (written <= nbytes) {
            bufp += written;
            nbytes -= written;
        }
    }
}

int main(int argc, char *argv[])
{
    char *source;
    char *dest;
    int source_fd, dest_fd;
    char buf[1024];
    ssize_t read_len;

    if (argc != 3) {
        printf("Usage: mycp <src-filename> <dest-filename>\n");
        exit(1);
    }

    source = argv[1];
    dest = argv[2];

    source_fd = open(source, O_RDONLY);
    if (source_fd == -1)
        error(1, errno, "error opening file %s", source);

    dest_fd = open(dest, O_WRONLY);
    if (dest_fd == -1)
        error(1, errno, "error opening file %s", dest);

    while (1) {
        read_len = read(source_fd, buf, sizeof(buf));
        if (read_len == -1)
            error(1, errno, "error reading file %s", source);

        if (read_len == 0)
            break;

        write_buf(dest_fd, buf, read_len, dest);
    }

    if (close(source_fd) == -1) {
        error(1, errno, "error closing file %s", source);
    }

    if (close(dest_fd) == -1) {
        error(1, errno, "error closing file %s", dest);
    }

    return 0;
}
```

```
}
```

cp with creation and truncation.

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <fcntl.h>

#include <stdio.h>
#include <stdlib.h>
#include <error.h>
#include <errno.h>

void write_buf(int fd, char *bufp, size_t nbytes, char *filename)
{
    int written;

    while (nbytes != 0) {
        written = write(fd, bufp, nbytes);
        if (written == -1)
            error(1, errno, "error writing to file %s", filename);

        if (written <= nbytes) {
            bufp += written;
            nbytes -= written;
        }
    }
}

int main(int argc, char *argv[])
{
    char *source;
    char *dest;
    int source_fd, dest_fd;
    char buf[1024];
    ssize_t read_len;
    struct stat stat;
    int ret;

    if (argc != 3) {
        printf("Usage: mycp <src-filename> <dest-filename>\n");
        exit(1);
    }

    source = argv[1];
    dest = argv[2];

    source_fd = open(source, O_RDONLY);
    if (source_fd == -1)
        error(1, errno, "error opening file %s", source);

    ret = fstat(source_fd, &stat);
```

```
    if (ret == -1)
        error(1, errno, "error getting mode bits of %s", source);

    dest_fd = open(dest, O_WRONLY | O_TRUNC | O_CREAT, /* ❶ */
                  stat.st_mode & 0777); /* ❷ */
    if (dest_fd == -1)
        error(1, errno, "error opening file %s", dest);

    while (1) {
        read_len = read(source_fd, buf, sizeof(buf));
        if (read_len == -1)
            error(1, errno, "error reading file %s", source);

        if (read_len == 0)
            break;

        write_buf(dest_fd, buf, read_len, dest);
    }

    if (close(source_fd) == -1) {
        error(1, errno, "error closing file %s", source);
    }

    if (close(dest_fd) == -1) {
        error(1, errno, "error closing file %s", dest);
    }

    return 0;
}
```

- ❶ Create file if not present. Truncate file if already present.
- ❷ Mode to use when file is created.